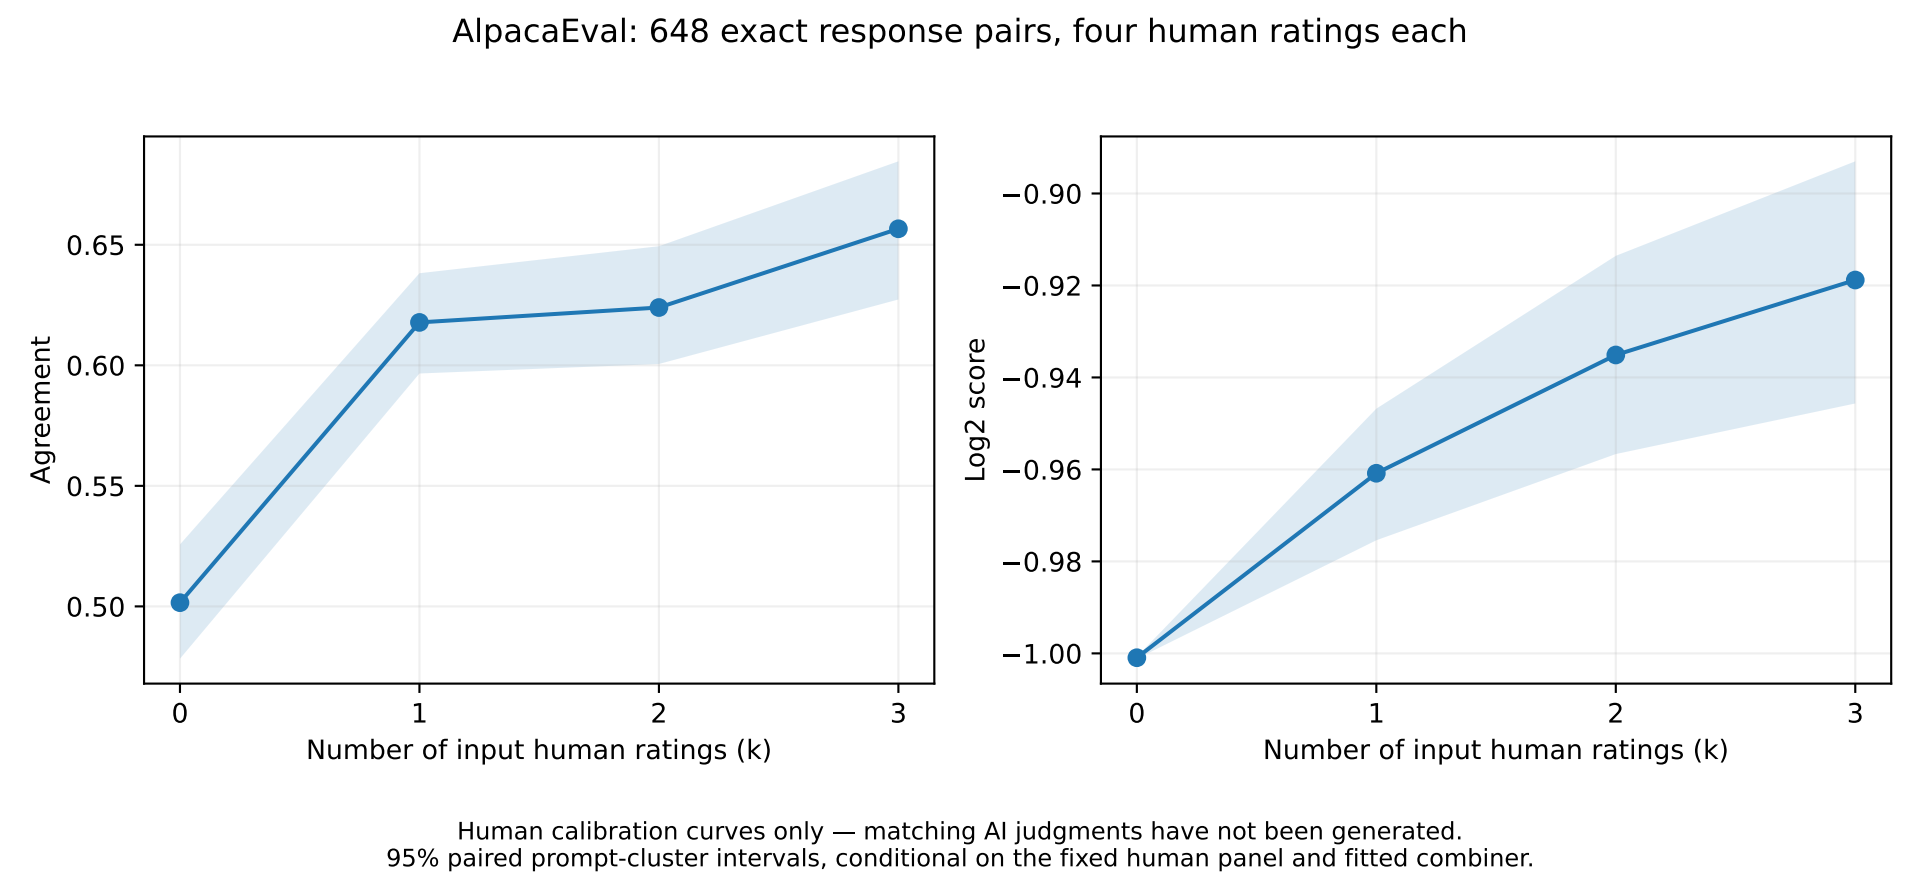
AlpacaEval: 648 exact response pairs, four human ratings each
Agreement
0.65
0.60
0.55
0.50
0
1
2
3
Number of input human ratings (k)
Log2 score
−0.90
−0.92
−0.94
−0.96
−0.98
−1.00
0
1
2
3
Number of input human ratings (k)
Human calibration curves only — matching AI judgments have not been generated.
95% paired prompt-cluster intervals, conditional on the fixed human panel and fitted combiner.