Arena: how well do AI judges predict human preferences?

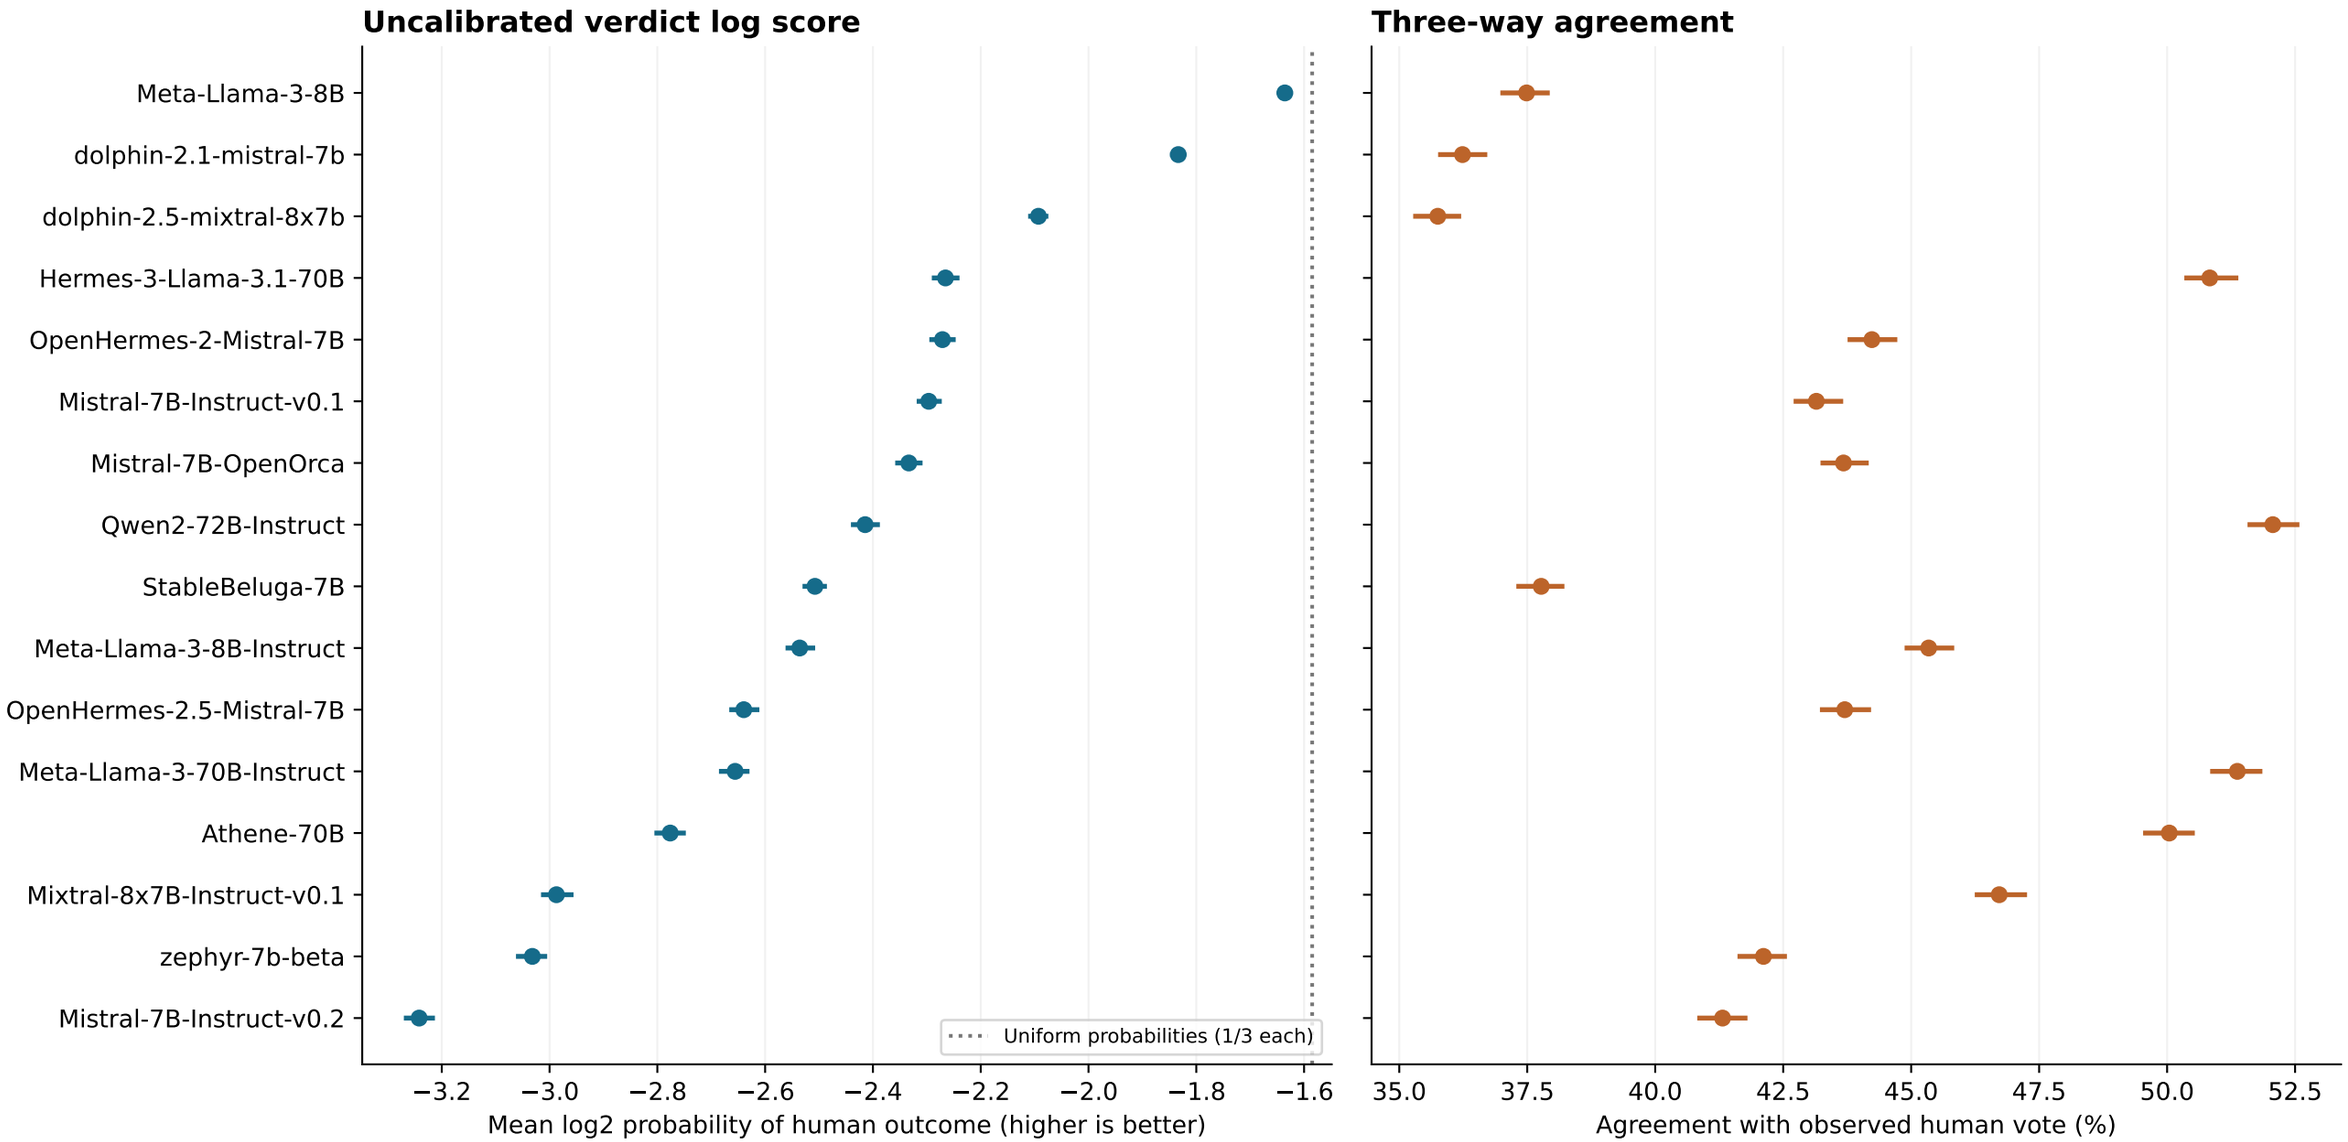

49,938 matched real battles · 16 verified judges · 95% intervals from 500 paired prompt-cluster bootstraps
Forward presentation; probabilities clipped at .02 and renormalized. One human vote per battle: human-panel equivalence is not identifiable.
Excluded for a published-baseline mismatch: Starling-LM-7B-alpha. Hard agreement uses the released strict-winner/tie rule.