MT-Bench: exploratory human-panel equivalence

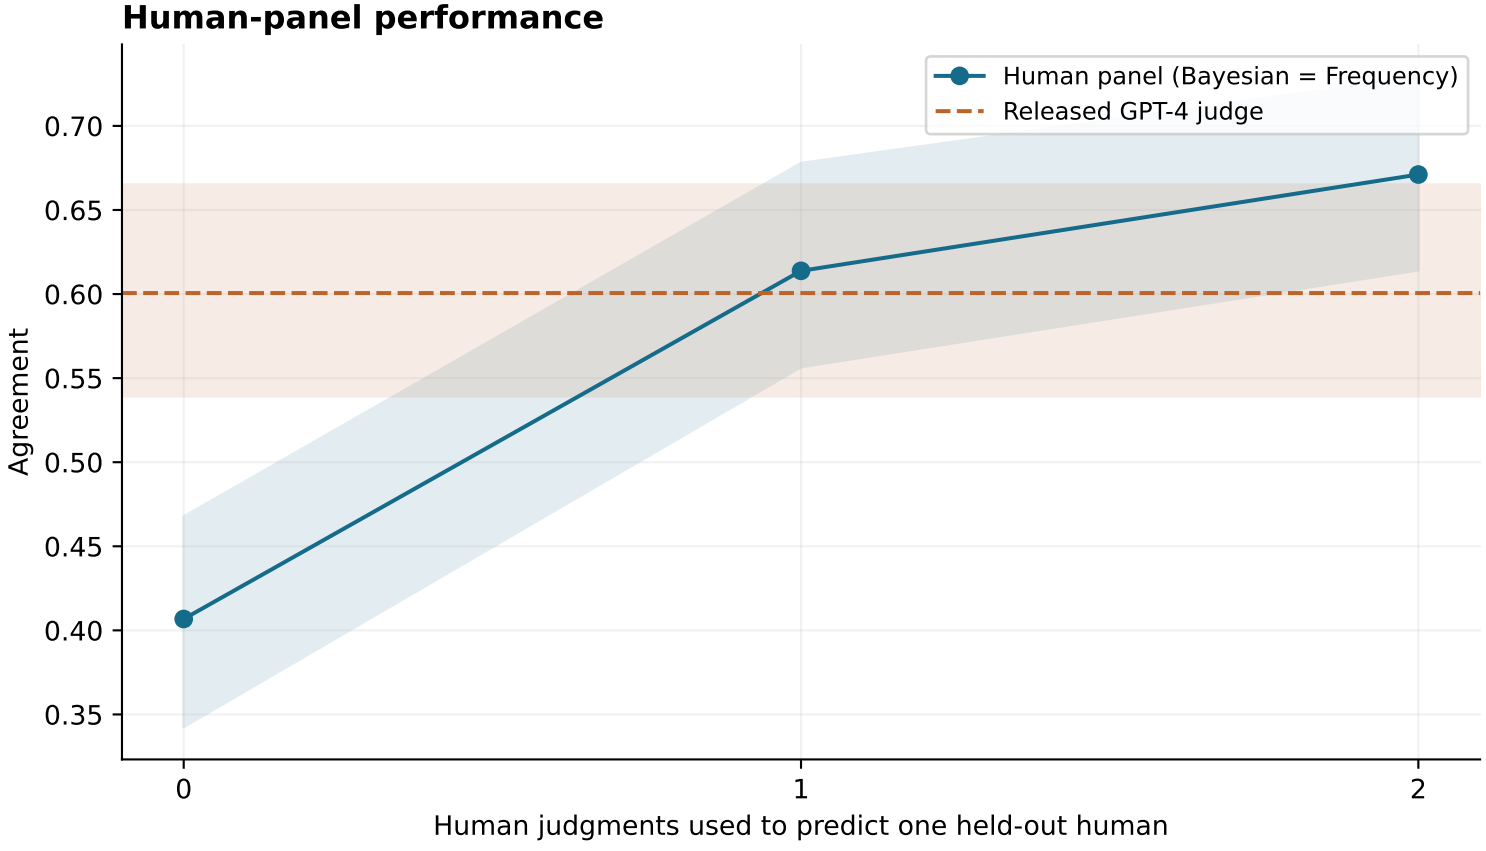

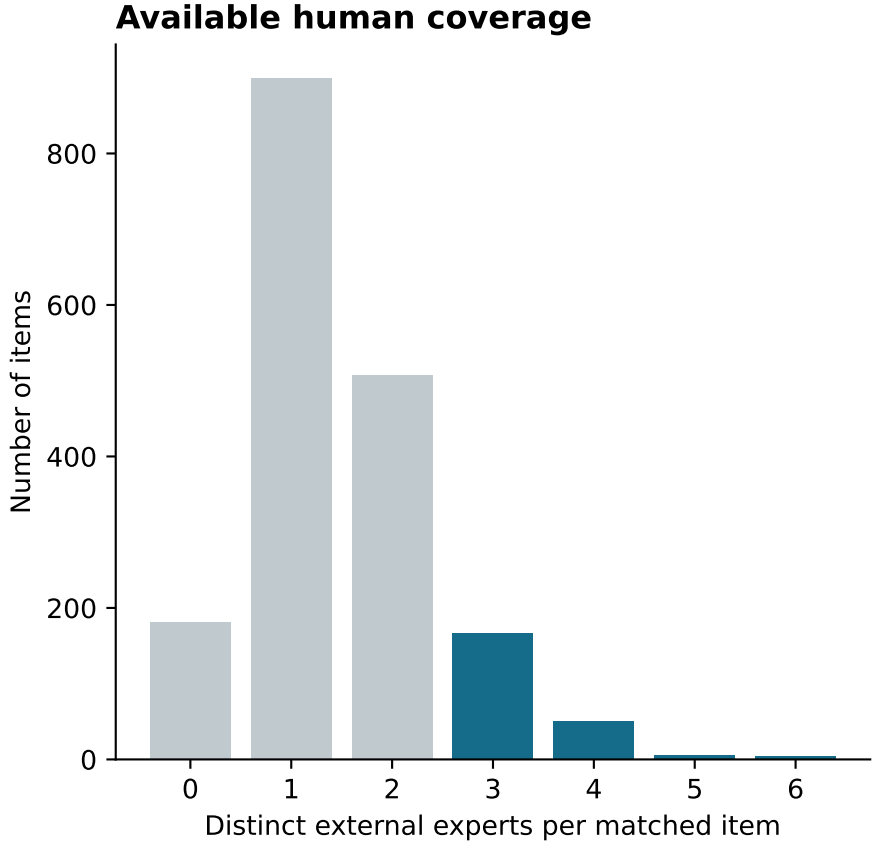

227 exact items · 3 distinct experts per item · GPT-4 Bayesian-reference equivalence: 0.94 humans (95% interval 0.68–1.70)
500 question-cluster bootstrap replicates; hard judgments only. Released vote independence from GPT-4 feedback is unverified.
Frequency k=0 uses its full-data prior. This cohort and GPT-4 judge differ from the Arena comparison above.