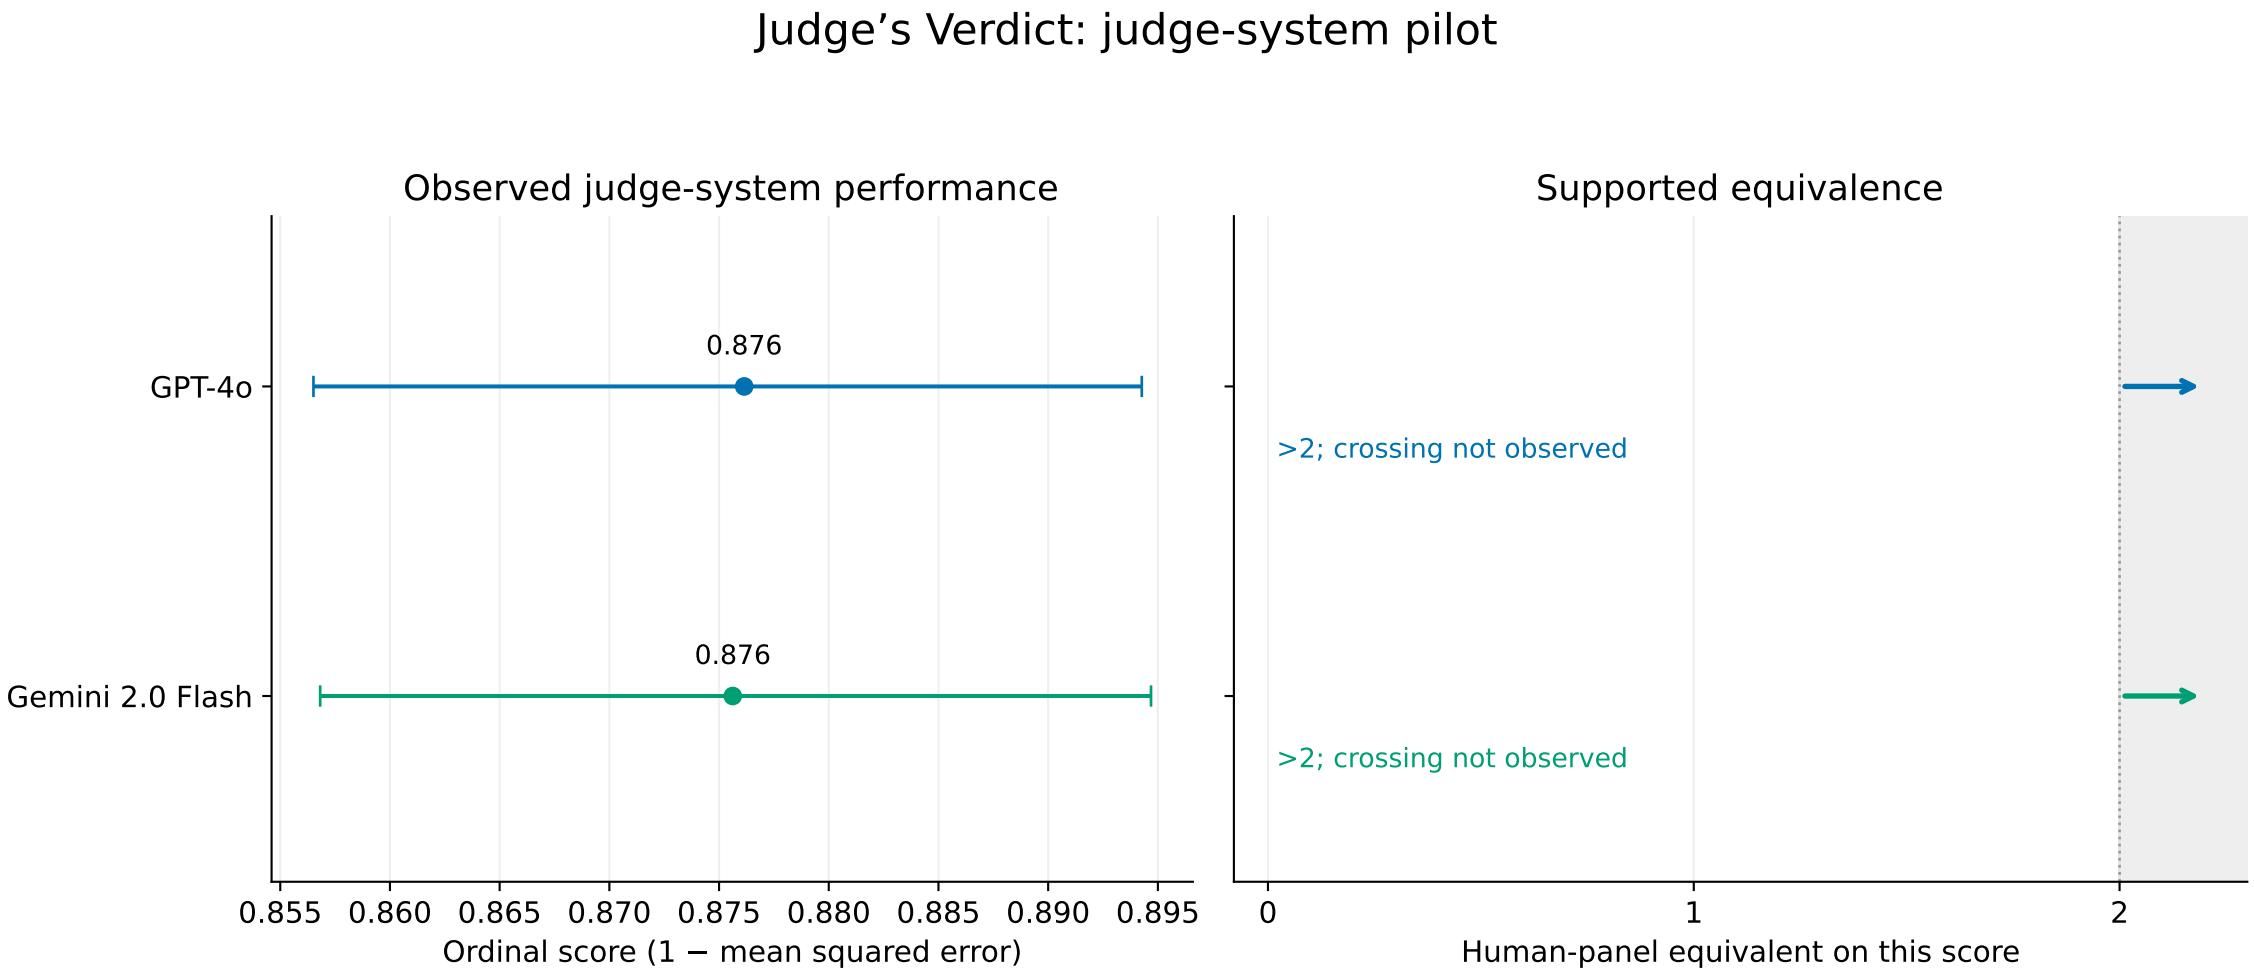

200 TechQA items; 3 human ratings/item; released trial 1 (two-prompt AI systems).
95% question-bootstrap intervals, conditional on fixed ratings and fitted combiner. Exploratory train subset.
Custom ordinal-score equivalence; not log-score / information equivalence.
Ordering by point score does not establish a statistically reliable ranking.